

1. Github

Github's GDPR guidelines can be roughly divided into three categories, necessary processing, legal compliance and consent-based actions [1]. Being able to log in is a necessity, so this goes in line with Art 6 (1) b) GDPR [2]. Github, like other companies, has a legal obligation to respond to data subject requests and disputes etc., which is handled by Art 6 (1) c) GDPR [2]. The third category, marketing and other non-essential interests are based fully on user consent, which is required by Art 6 (1) a) GDPR [2]. If the user wishes to withdraw their consent or remove their personal data, they have the legal right to do so, as stated in Art 7 (3) GDPR and Art 17 (1) GDPR. Github claims to follow this law in their privacy statement [1]. Based on Github's own claims, they seem to be following GDPR guidelines.

Sources:

[1] <https://docs.github.com/en/site-policy/privacy-policies/github-general-privacy-statement>

[2] <https://gdpr-info.eu/art-6-gdpr/>

2. An individual's rights as a data subject are defined in Art 15 GDPR [1]. These rights include being able to know what personal data the service has (1) a), how long this data will be stored for (1) d), and requests for erasure of personal data (1) e) etc. According to Github's privacy statement [2], users can exercise these rights by sending an email to privacy@github.com. Many of the rights do include clauses like "under certain circumstances" and "where technically feasible", which does leave me doubting the effectiveness of this system. It seems that even though GDPR laws exist and they are mostly effective, companies do try to find ways around these systems. Personally however, I did find my answer from Github satisfactory enough.

Sources:

[1] <https://gdpr-info.eu/art-15-gdpr/>

[2] <https://docs.github.com/en/site-policy/privacy-policies/github-general-privacy-statement>

3. Art 4 (5) GDPR defines pseudonymization as processing personal data in a way where it's not possible to tell whose data it was originally. Additional data may be kept elsewhere that can be used to attribute this data to a specific person. I was not able to find the corresponding definition of anonymization in the GDPR, but in practice it refers to irreversibly removing information that could attribute that data to a specific person. Practical example of pseudonymization:

DB1:

John Smith -> ID0004

DB2:

ID0004 -> John Smith

In anonymization there would be no DB2.

Exercise 2

1. Reducing the salary of women has practically no effect on the slope. Changing the binomial distribution parameters also seems to have no effect on the slope. Changing both aspects seems to affect the slope a little more, but it doesn't seem to have a major effect. Overall, none of the three seem to do much.
2. To detect indirect discrimination the protected attribute Gender should be included. This could also be compared to the Hours statistic. If we fit both gender and hours into linear regression and print out the coefficient, we can clearly see discrimination between the two genders e.g [-198, 100.9].